

Introduction

INF 551

Wensheng Wu

Logistics

- Instructor email: wenshenw@usc.edu
- Class meeting times:
 - Morning section: 10-11:50am, MW, VPD 106
 - Afternoon section: 3:30-5:20pm, MW, THH 118
- Office hours:
 - 9-9:45am, MW, GER 204 (please email me first to set up a time)

Logistics

- TAs
 - Weijun Deng, weijunde@usc.edu (morning)
 - Siyan Cheng, siyanc@usc.edu (afternoon)
- Office hours
 - TBD
 - SAL computing lab (lobby)

Blackboard

- Discussion forums
 - You may post general and homework questions
 - Do not post solutions
 - Please actively participate in helping others!
 - Do not abuse forum (an academic misconduct!)
- Check frequently for updates

Prerequisites

- Programming skills:
 - Python (e.g., for Spark), Java (e.g., for Hadoop)
- Unix-like environment
 - E.g., Ubuntu, Virtual machine, Amazon EC2 (we will use this)
- Basic knowledge of algorithms and data structures
 - Sorting, hashing, etc. (CS 570)
- Basic probability and statistics

Textbooks

- Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau. *Operating Systems: Three Easy Pieces*, 2015 (selected chapters only). Available free at:
<http://pages.cs.wisc.edu/~remzi/OSTEP/>
- Hector Garcia-Molina, Jeffrey D. Ullman, and Jennifer Widom. *Database Systems: The Complete Book* (Second Edition), Prentice Hall, 2009. (selected chapters only)
 - <http://infolab.stanford.edu/~ullman/dscb.html>
- Jiawei Han, Micheline Kamber, and Jian Pei. [Data Mining: Concepts and Techniques](#). Morgan Kaufmann, 2011, 3rd Edition (selected chapters only).

Additional readings

- Posted on Blackboard
- Base your reviews on the lectures
 - I will not ask questions on the reading materials that were not covered in the lectures

Grading structure

- Homework 20%
- Weekly quizzes 25%
- Midterm 15%
- Final 25%
- Lab session 5%
- Group project 10%

Grading scale

- $[93, 100] = A$
- $[90, 93) = A-$
- $[87, 90) = B+$
- $[83, 87) = B$
- $[80, 83) = B-$
- $[77, 80) = C+$
- $[73, 77) = C$
- ... (see Syllabus for complete breakdown)

Lab sessions

- Flipped:
 - Task and details posted before class
 - Bring questions to class
- Typically utilize last 15-30 mins of class

Quizzes

- First quiz on 3rd week
- Last quiz on last week
- Based on previous week's materials

Exams

- Closed-notes & book
- Midterm: Oct. 3, Wednesday, in-class
- Final: December 10, Wednesday
 - Morning: 8-10am; afternoon section: 2-4pm
 - At the same classroom as class meeting

Calculator

- If calculator is needed, we will announce it
- Otherwise, no electronic devices are allowed

Group project

- Form a group of no more than 2 people
- Done in phases
 - Proposal
 - Midterm report
 - Final report
- Will talk about project topics

Late Policy

- Homework will be submitted to Blackboard
 - 10% for every 24 hours late
 - No credit after 3 days
- Make up for quizzes are permitted only when
 - You have a medical/family emergency
 - Let me know in advance
 - Proof (e.g., medical note) is required

Late Policy

- Quiz will be given in the beginning of Monday's classes
 - You are responsible for missing quiz due to tardiness
 - No make up will be given for tardiness!
- You are responsible for scheduling conflicts
 - With job interviews, job fairs, etc.

Grading Corrections

- All homework & quiz grades are final one week after grades are posted
- Wasting instructor and staff times on non-legitimate regrading requests may result in losing more points

Academic Integrity

- **Cheating will NOT be tolerated**
- **All parties involved will receive a grade of F for the course and be reported to SJACS WITHOUT EXCEPTION**
 - [USC Student Judicial Affairs and Community Standards](#)


Now, movie time 😊

- Explain big data:
 - https://www.youtube.com/watch?v=7D1CQ_LOizA
- Questions:
 - Where does big data come from?
 - What characteristics does it have?
 - What big data technologies were mentioned?

Internet Traffic in 2012

- 4.8 zettabyte = 4.8 billion terabytes
- Zettabyte (1000 exabytes)
- Exabyte
- Petabyte
- Terabyte
- Gigabyte
- Megabyte
- Kilobyte

Major topics

- Storage systems 
- File systems & file formats
- Database management systems
- Big data solution stack
- Data warehousing (if time permits)

Storage Systems

- Hard disk
- SSD (Solid state drive)



Internal of hard disk

Actuator


Spindle



Platter

Disk head

NAND flash



Major topics

- Storage systems
- **File systems** & file formats
- Database management systems
- Big data solution stack
- Data warehousing



File Systems


- Standalone
 - Single machine
- Network
 - Client-server
- Distributed (e.g., Hadoop)
 - A number of data servers

Standalone file systems

- Data structures
 - Data blocks
 - Metadata blocks (Inodes)
 - Bitmap blocks (for space allocation)
- Access paths
 - Read
 - write


Inode (index node)

- Each is identified by a number
 - Low-level number of file name: inumber
- Can figure out location of inode from inumber



Network file system

- Client-server architecture
 - Sun's network file system
- Key concept:
 - stateless file handle



Distributed file systems

- Hadoop HDFS (after GFS)
 - Data are distributed among data nodes
- Replication
 - Automatic creation of replica (typically 2 or 3 copies/data)
- Fault-tolerant
 - Automatic recovery from node failure

HDFS architecture

NameNode:


Stores metadata only

METADATA:

/user/aaron/foo → 1, 2, 4


/user/aaron/bar → 3, 5

DataNodes: Store blocks from files



Major topics

- Storage systems
- File systems & **file formats**
- Database management systems
- Data warehousing
- Big data solution stack



File Formats

- JSON

```
{
  "firstName": "John",
  "lastName": "Smith",
  "isAlive": true,
  "age": 25,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null
}
```

HTML

<h1> Bibliography </h1>

<p> <i> Foundations of Databases </i>

Abiteboul, Hull, Vianu

 Addison Wesley, 1995

<p> <i> Data on the Web </i>

Abiteoul, Buneman, Suciu

 Morgan Kaufmann, 1999

XML

```
<bibliography>
  <book>  <title> Foundations... </title>
          <author> Abiteboul </author>
          <author> Hull </author>
          <author> Vianu </author>
          <publisher> Addison Wesley </publisher>
          <year> 1995 </year>
  </book>
  ...
</bibliography>
```

XML describes the content

XML usages

- Software configurations files
 - E.g., HDFS
- Android app development
 - Layout resource files, e.g., activity_main.xml
- Java archive (.jar file)
 - Manifest.xml

Android app resource file

```
<?xml version="1.0" encoding="utf-8"?>
<RelativeLayout xmlns:android="http://schemas.android.com/apk/res/android"
    xmlns:tools="http://schemas.android.com/tools"
    android:layout_width="match_parent"
    android:layout_height="match_parent"
    tools:context=".MainActivity">

    <android.support.design.widget.TabLayout
        android:id="@+id/tabs"
        android:layout_width="match_parent"
        android:layout_height="wrap_content" />

    <android.support.v4.view.ViewPager
        android:id="@+id/container"
        android:layout_width="match_parent"
        android:layout_height="match_parent"
        android:layout_below="@id/tabs" />
```

Manifest.xml

```
<?xml version="1.0" encoding="utf-8"?>
<manifest xmlns:android="http://schemas.android.com/apk/res/android"
    package="com.google.firebase.quickstart.database">


    <uses-permission android:name="android.permission.INTERNET" />

    <application
        android:allowBackup="true"
        android:icon="@mipmap/ic_launcher"
        android:label="Firebase Database"
        android:supportsRtl="true"
        android:theme="@style/AppTheme">
        <activity
            android:name=".MainActivity"
            android:label="Firebase Database"
            android:theme="@style/AppTheme" />
        <activity android:name=".NewPostActivity" />
        <activity android:name=".SignInActivity">
            <intent-filter>
                <action android:name="android.intent.action.MAIN" />

                <category android:name="android.intent.category.LAUNCHER" />
            </intent-filter>
        </activity>
    </application>
</manifest>
```

Major topics

- Storage systems
- File systems & file formats
- Database management systems
- Big data solution stack
- Data warehousing



Relational DBMS

- Data models
 - ER
 - Relational
- Schema
 - Normal forms: BCNF

RDBMS

- Query languages
 - Relational algebra
 - SQL, constraints, views
- Data organization
 - Records and blocks
 - Index structure: B+-tree

RDBMS

- Query execution algorithms
 - External sorting
 - One-pass algorithms
 - Nested-loop join
 - Multiple-pass algorithms

RDBMS

- Rigid schema
- Strong consistency is the key design goal
 - Never read old data
 - Suitable for mission-critical applications, e.g., banking
- But may suffer from low availability

RDBMS







- Hard to scale out
 - Horizontal partitioning/sharding possible
 - But would need distributed storage & computing support like Hadoop & MapReduce

RDBMS Examples

- MySQL (can be installed in Amazon AWS EC2)
- Amazon RDS (Relational database service)
 - DBMS in the cloud
 - Database as a service
- Data warehouse on RDBMS
 - OLAP

Amazon RDS: Database-as-a-service

- MySQL, PostgreSQL, Oracle, SQL Server, etc.

	<h2>MySQL</h2> <p>MySQL Community Edition</p> <p>MySQL is the most popular open source database in the world. MySQL on RDS offers the rich features of the MySQL community edition with the flexibility to easily scale compute resources or storage capacity for your database.</p> <ul style="list-style-type: none">• Supports database size up to 6 TB.• Instances offer up to 32 vCPUs and 244 GiB Memory.• Supports automated backup and point-in-time recovery.• Supports cross-region read replicas.	Select
		
		
		
		
		

Access MySQL from EC2

```
ssh -i "weixin.pem" ec2-user@ec2-107-22-155-60.compute-1.amazonaws.com
Last login: Sat Aug 6 07:43:38 2016 from 211.162.33.156

 _ _ | _ _ | _ _ )
 _ _ | ( _ _ | _ _ /
 _ _ | \ _ _ | _ _ |

Amazon Linux AMI

https://aws.amazon.com/amazon-linux-ami/2016.03-release-notes/
[ec2-user@ip-172-31-50-20 ~]$ mysql
ERROR 2002 (HY000): Can't connect to local MySQL server through socket '
/var/lib/mysql/mysql.sock' (2)
[ec2-user@ip-172-31-50-20 ~]$ sudo service mysqld start
Starting mysqld: [ OK ]
[ec2-user@ip-172-31-50-20 ~]$ mysql
ERROR 1045 (28000): Access denied for user 'ec2-user'@'localhost' (using
password: NO)
[ec2-user@ip-172-31-50-20 ~]$ mysql -h inf551.chdcdeeoqxf5.us-east-1.rds
.amazonaws.com -P 3306 -u inf551 -p
Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 111
Server version: 5.6.27-log MySQL Community Server (GPL)

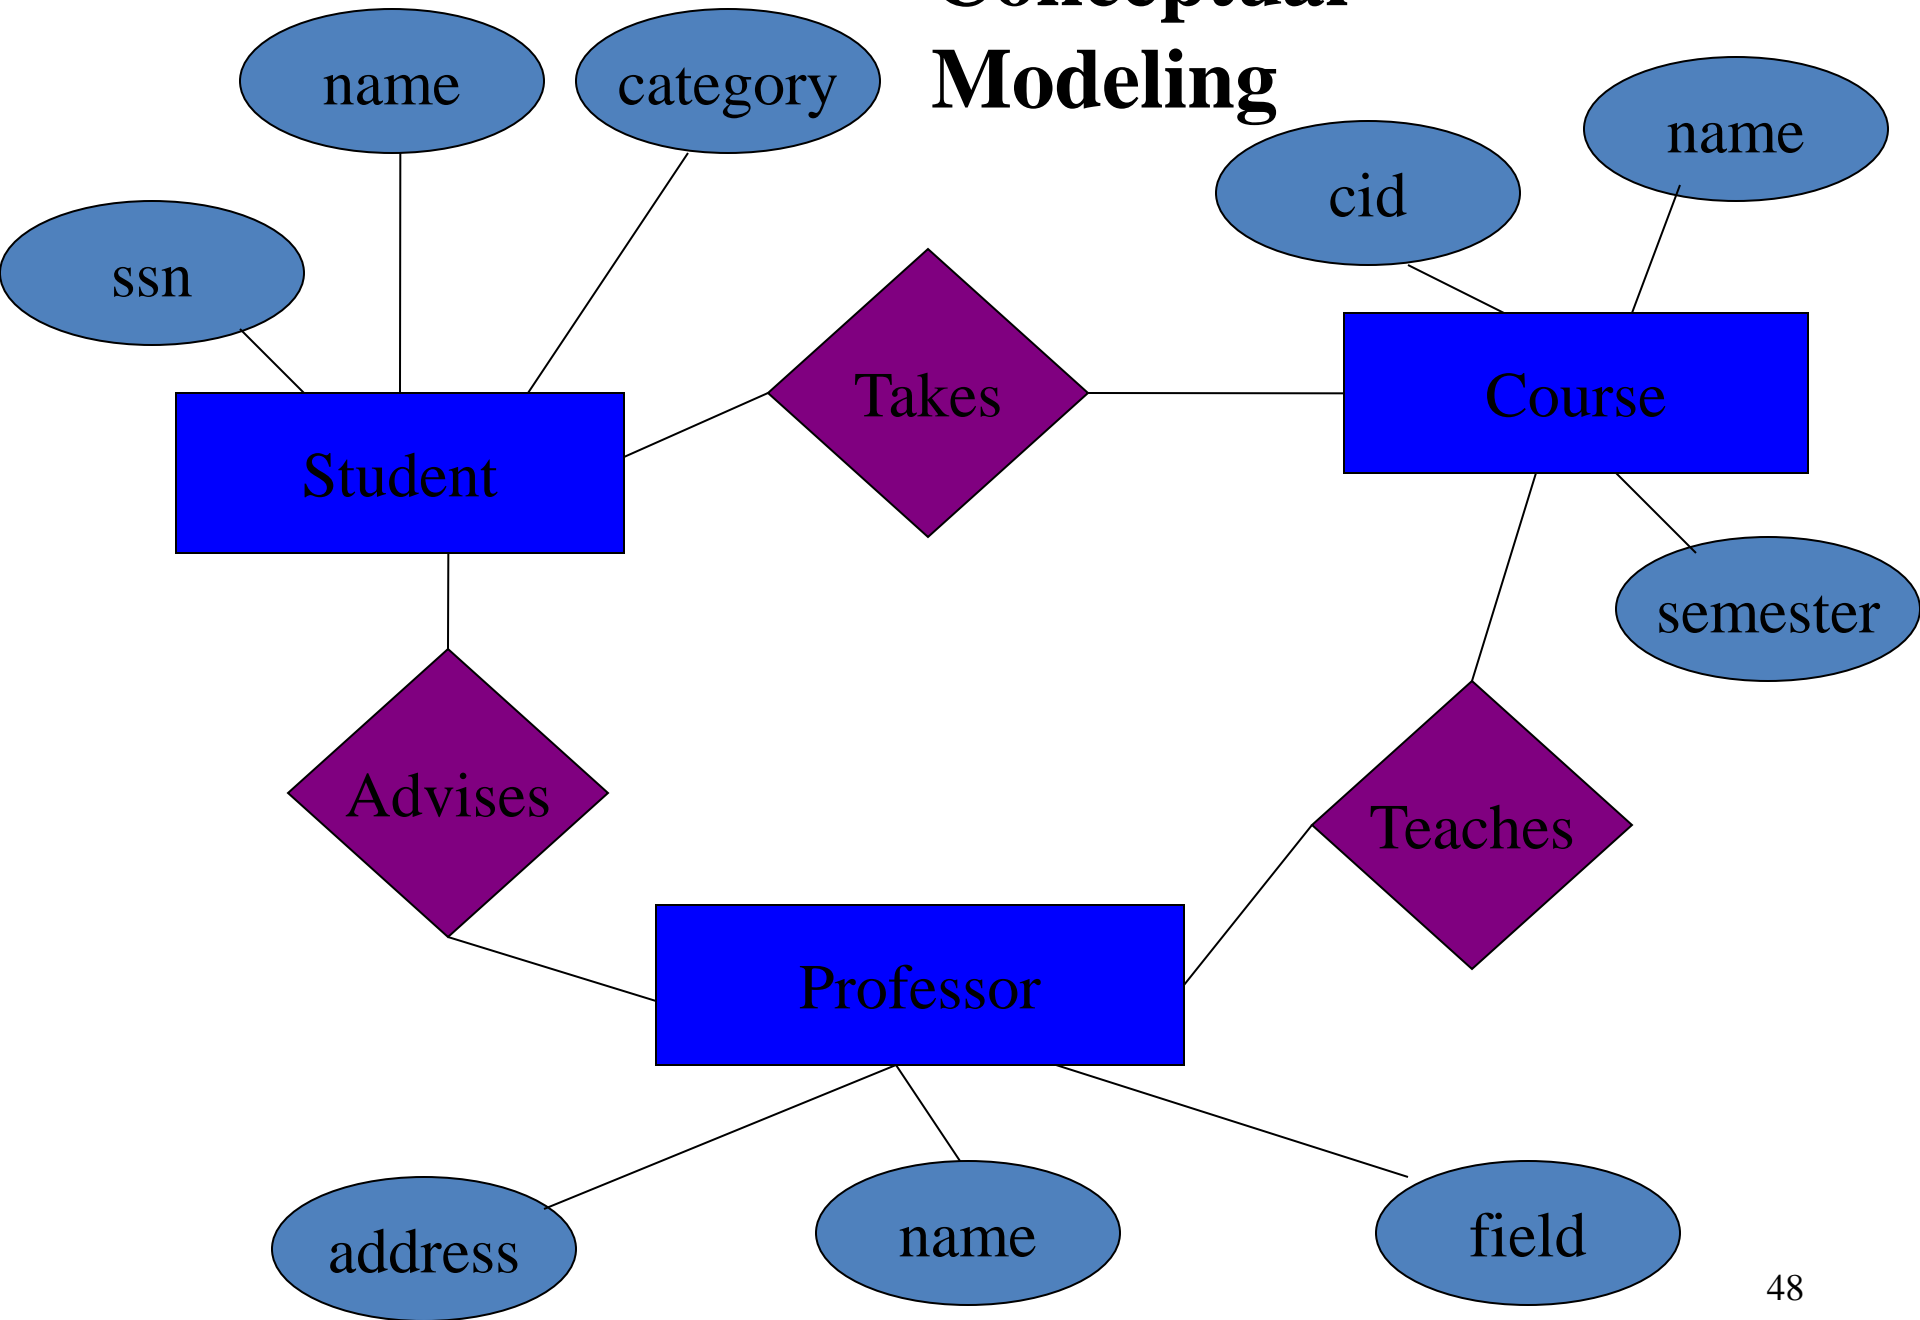
Copyright (c) 2000, 2015, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql>
```

Conceptual Modeling



Schema Design and Implementation

- Tables:

Students:

SSN	Name	Category
123-45-6789	Charles	undergrad
234-56-7890	Dan	grad

Takes:

SSN	CID
123-45-6789	CSE444
123-45-6789	CSE444
234-56-7890	CSE142
	...

Courses:

CID	Name	Semster
CSE444	Databases	fall
CSE541	Operating systems	spring

- Separates the logical view from the physical view of the data.

Querying a Database

- Find all courses that "Mary" takes
- S(tructured) Q(uey) L(anguage)

```
select C.name  
from    Students S, Takes T, Courses C  
where   S.name = "Mary" and  
         S.ssn = T.ssn and T.cid = C.cid
```

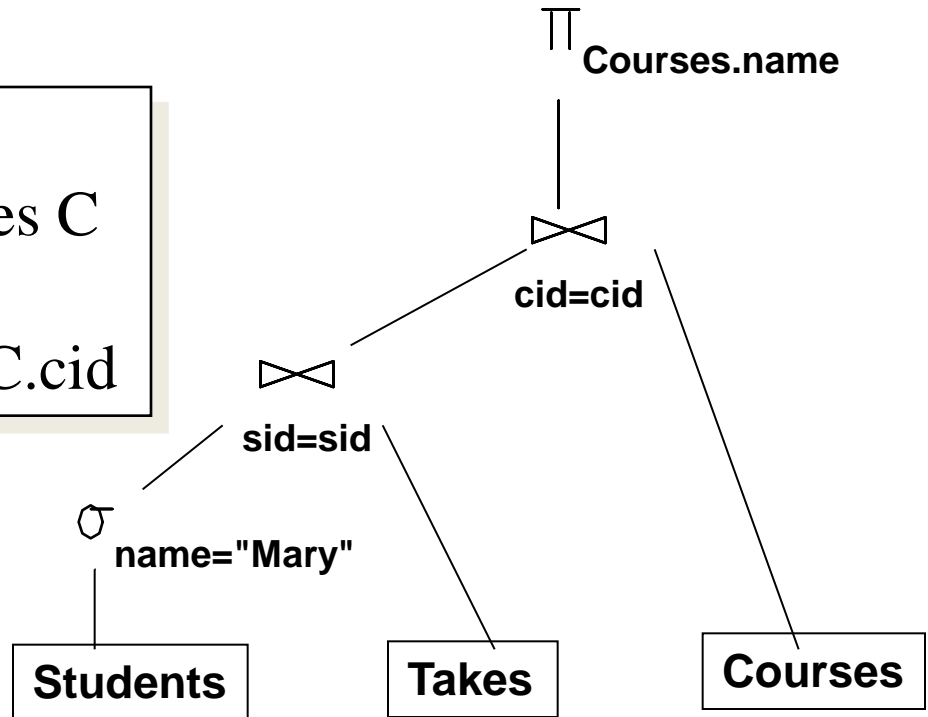
- Query processor figures out how to answer the query efficiently.

Query Optimization

Goal:


Declarative SQL query • \longrightarrow *Imperative query execution plan:*

```
select C.name
from Students S, Takes T, Courses C
where S.name="Mary" and
      S.ssn = T.ssn and T.cid = C.cid
```



Plan: tree of Relational Algebra operators,
choice of algorithms at each operator

Major topics

- Storage systems
- File systems & file formats
- Database management systems
- Big data solution stack 
- Data warehousing

Topics

- Big data management & analytics
 - Cloud data storage (Amazon S3)
 - NoSQL (Amazon DynamoDB, Cassandra, MongoDB, Google Firebase)
 - MapReduce
 - Apache Hadoop
 - Apache Spark
 - Apache Hive (if time permits)

Cloud data storage

- Amazon S3 (simple storage service)
 - Ideal for storing large binary files
 - E.g., audio, video, image
 - Simple RESTful web service
- Eventual consistency for high availability

PRODUCTS & SERVICES

- [Amazon S3](#) >
- [Product Details](#) >
- [Storage Classes](#) >
- [Pricing](#) >
- [Getting Started](#) >
- [FAQs](#) >
- [Resources](#) >
- [Amazon S3 SLA](#) >

RELATED LINKS

- [AWS Management Console](#)
- [Documentation](#)
- [Release Notes](#)

Amazon S3

Amazon Simple Storage Service (Amazon S3), provides developers and IT teams with secure, durable, highly-scalable object storage. Amazon S3 is easy to use, with a simple web service interface to store and retrieve any amount of data from anywhere on the web. With Amazon S3, you pay only for the storage you actually use. There is no minimum fee and no setup cost.

Amazon S3 offers a range of storage classes designed for different use cases including Amazon S3 Standard for general-purpose storage of frequently accessed data, Amazon S3 Standard - Infrequent

Get Started with AWS Today

[Try Amazon S3 for Free](#)

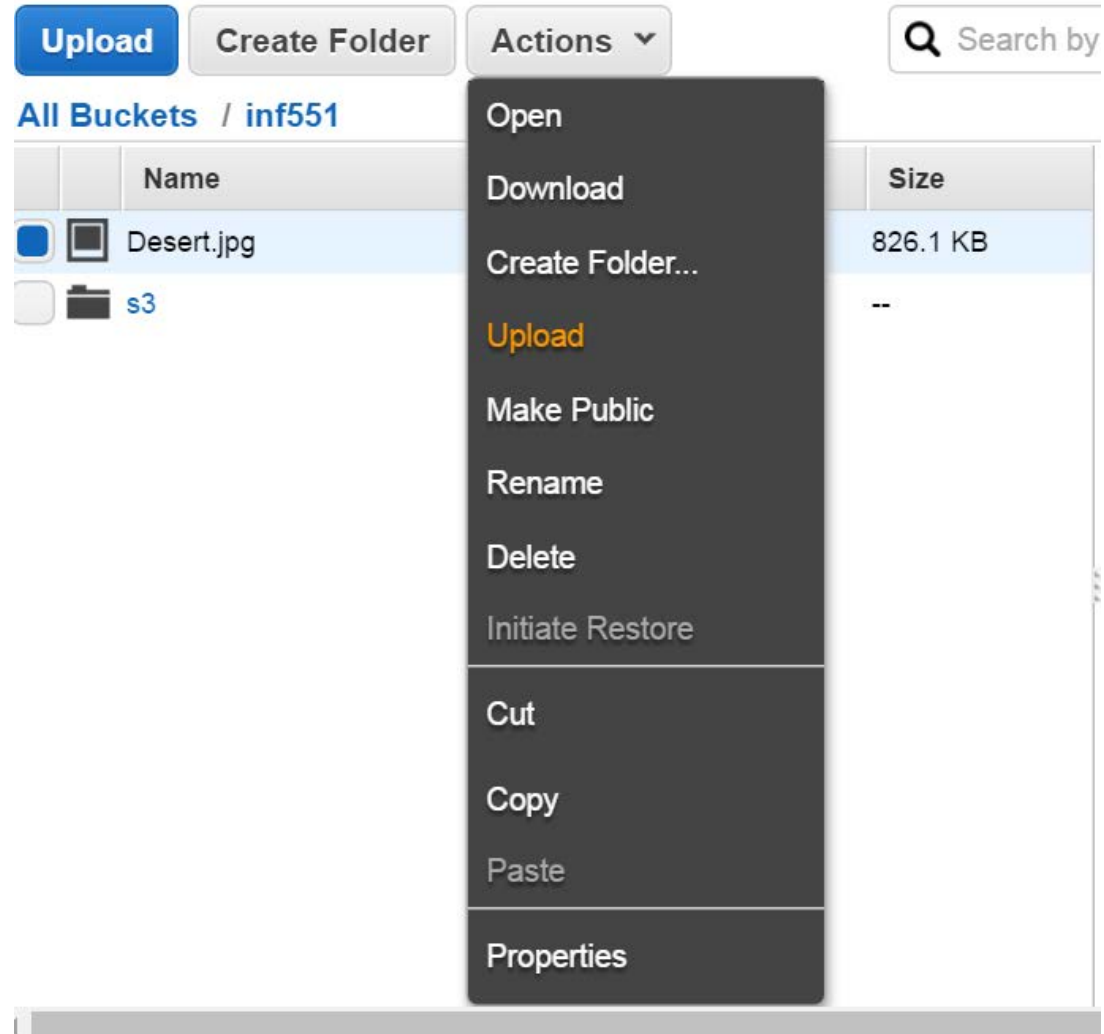
AWS Free Tier includes 5GB storage, 20,000 Get Requests, and 2,000 Put Requests with Amazon S3.

[View AWS Free Tier Details »](#)

In Recent News

New: Amazon VPC

Upload a file





<https://s3.amazonaws.com/inf551/Desert.jpg>



NoSQL

- Not only SQL
- Flexible schemas
 - e.g., JSON documents or key-value pairs
 - Ideal for managing a mix of structured, semi-structured, and unstructured data
- High availability
- Weaker (e.g., eventual) consistency model

Example NoSQL databases

- Amazon DynamoDB
 - Row store
 - row = item = a collection of key-value pairs
- MongoDB, Firebase, etc.
 - Manage JSON documents
- Apache Cassandra
 - Wide column store
 - Google's Bigtable clone

Key techniques

- Consistent hashing (Cassandra, Dynamo)
 - Avoid moving too much data when adding new machines
- Efficient writes
 - Append-only
 - No overwrites
 - Avoid random seek
 - But compaction needed later

Key techniques

- Compaction
 - Introduced in Google "Bigtable" paper
 - Merge multiple versions of data
 - Remove expired or deleted data

DynamoDB

- <https://console.aws.amazon.com/dynamodb/home?region=us-east-1#gettingStarted>:

Amazon DynamoDB

Amazon DynamoDB is a fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale. Its flexible data model and reliable performance make it a great fit for mobile, web, gaming, ad-tech, IoT, and many other applications.

Create table

Create DynamoDB table

Tutorial



DynamoDB is a schema-less database that only requires a table name and primary key. The table's primary key is made up of one or two attributes that uniquely identify items, partition the data, and sort data within each partition.

Table name*

Books



Primary key*

Partition key

Author

String



☒ Add sort key

Title

String



Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

☒ Use default settings

Insert items

DynamoDB

Dashboard

Tables

Reserved capacity

Create tableActions ▾

Filter by table name X

	Name
<input checked="" type="radio"/>	Books

OverviewItemsMetricsAlarms

Create itemActions ▾

Scan: [Table] Books: Author, Title ^

Scan ▾

[Table] Books: Author, Title

+ Add filter

Start search

<

<input type="checkbox"/>	Author	Title	
--------------------------	--------	-------	--

May add new attributes

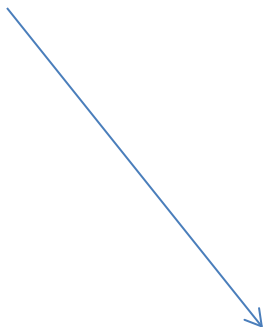
Tree ▾

Item {3}

- Author String : Bill Clinton
- Title String : My Life
- ISBN String : 1234567890

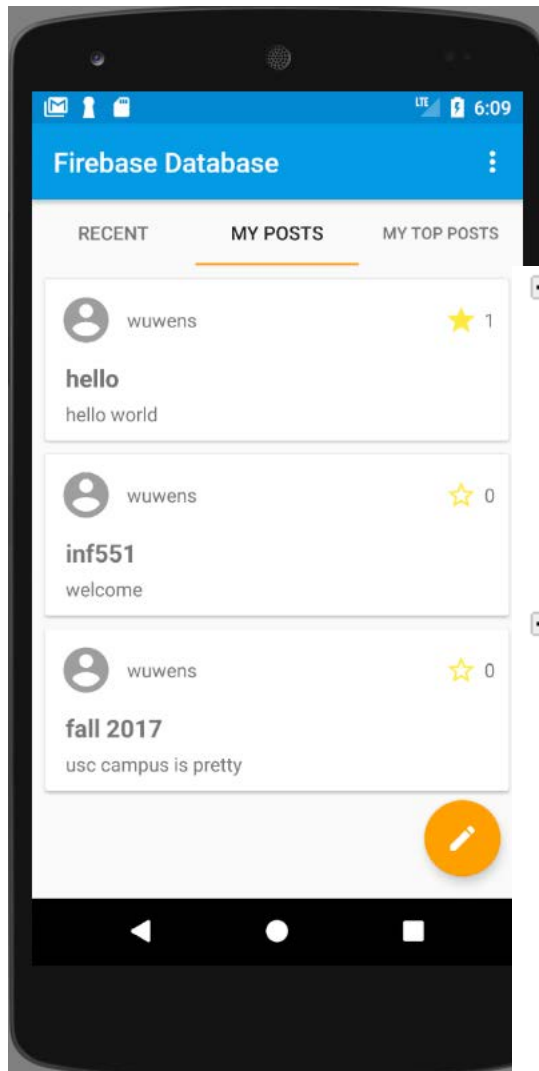
+ Append ▾

- String
- Binary
- Number
- StringSet
- NumberSet
- BinarySet
- Map
- List



<input type="checkbox"/>	Author	Title	ISBN
<input type="checkbox"/>	Bill Clinton	My Life	1234567890

Firestore: a cloud database



post-comments

-Ks-oZBfVttBkx_Ocfrq

-Ks-otimnHiahFzpzgvy

author: "wuwens"

text: "hello hello"

uid: "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"

posts

+ -Ks-oNK4isG3XjvO0ym2 + -

+ -Ks-oPc7HLjKN8Z5p9j7

-Ks-oZBfVttBkx_Ocfrq

author: "wuwens"

body: "hello world"

starCount: 1

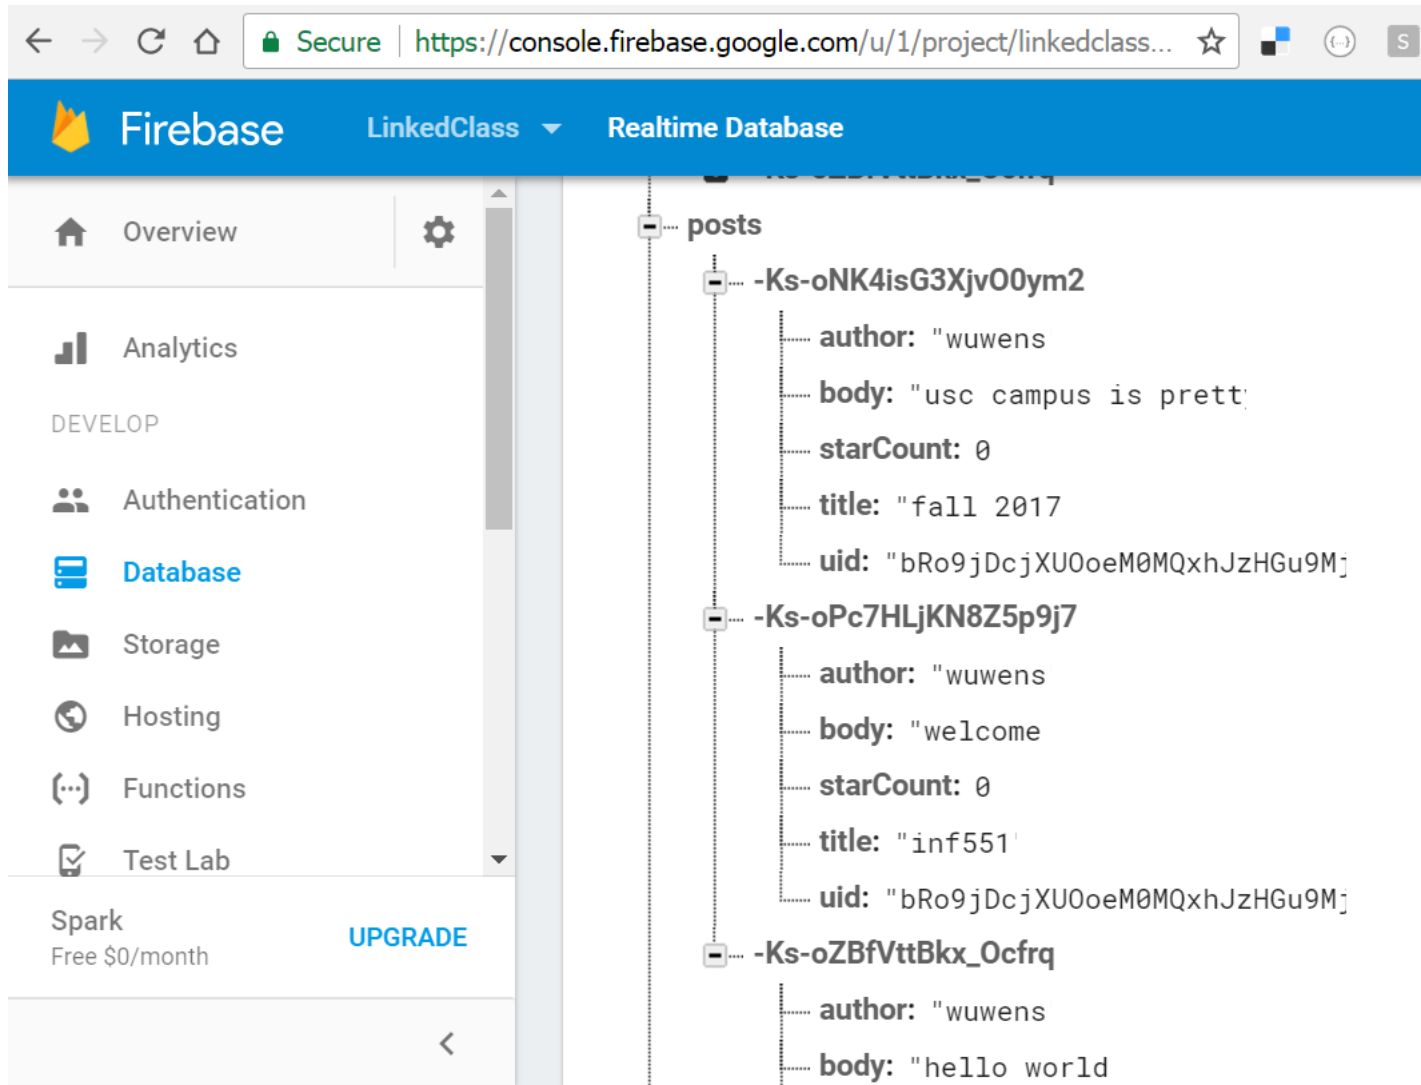
+ stars

title: "hello"

uid: "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"

```
{
  "post-comments" : {
    "-Ks-oZBfVttBkx_Ocfrq" : {
      "-Ks-otimnHiahFzpzgvy" : {
        "author" : "wuwens",
        "text" : "hello hello",
        "uid" : "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"
      }
    }
  },
  "posts" : {
    "-Ks-oNK4isG3XjvO0ym2" : {
      "author" : "wuwens",
      "body" : "usc campus is pretty",
      "starCount" : 0,
      "title" : "fall 2017",
      "uid" : "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"
    },
    "-Ks-oPc7HLjKN8Z5p9j7" : {
      "author" : "wuwens",
      "body" : "welcome",
      "starCount" : 0,
      "title" : "inf551",
      "uid" : "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"
    },
    "-Ks-oZBfVttBkx_Ocfrq" : {
      "author" : "wuwens",
      "body" : "hello world",
      "starCount" : 1,
      "stars" : {
        "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2" : true
      },
      "title" : "hello",
      "uid" : "bRo9jDcjXUOoeM0MQxhJzHGu9Mj2"
    }
  }
}
```

Firebase



The screenshot shows the Firebase Realtime Database console in a web browser. The address bar indicates the URL is <https://console.firebase.google.com/u/1/project/linkedclass...>. The page header includes the Firebase logo, the project name "LinkedClass", and the selected database "Realtime Database".


The left sidebar contains navigation links: Overview, Analytics, DEVELOP, Authentication, Database (highlighted), Storage, Hosting, Functions, and Test Lab. At the bottom of the sidebar, there is a "Spark" section indicating "Free \$0/month" and an "UPGRADE" button.

The main content area displays a tree view of the database structure. The "posts" node is expanded, showing three entries:

- posts**
 - Ks-oNK4isG3Xjv00ym2**
 - author: "wuwens"
 - body: "usc campus is prett"
 - starCount: 0
 - title: "fa11 2017"
 - uid: "bRo9jDcjXU0oeM0MQxhJzHGu9Mj"
 - Ks-oPc7HLjKN8Z5p9j7**
 - author: "wuwens"
 - body: "welcome"
 - starCount: 0
 - title: "inf551"
 - uid: "bRo9jDcjXU0oeM0MQxhJzHGu9Mj"
 - Ks-oZBfVttBkx_Ocfrq**
 - author: "wuwens"
 - body: "hello world"

Topics

- Big data management & analytics
 - Cloud data storage (Amazon S3)
 - NoSQL (Amazon DynamoDB, Cassandra, MongoDB)
 - **MapReduce**
 - Apache Hadoop
 - Apache Spark
 - Apache Hive



Roots in functional programming

- Functional programming languages:
 - Python, Lisp (list processor), Scheme, Erlang, Haskell
- Two functions:
 - Map: mapping a list => list
 - Reduce: reducing a list => value
- map() and reduce() in Python
 - <https://docs.python.org/2/library/functions.html#map>

map() and reduce() in Python

- `list = [1, 2, 3]`
- `def sqr(x): return x ** 2`
- `list1 = map(sqr, list)`

What are the value of list1 and z?



- `def add(x, y): return x + y`
- `z = reduce(add, list)`

Lambda function

- Anonymous function (not bound to a name)
- `list = [1, 2, 3]`
- `list1 = map(lambda x: x ** 2, list)`
- `z = reduce(lambda x, y: x + y, list)`

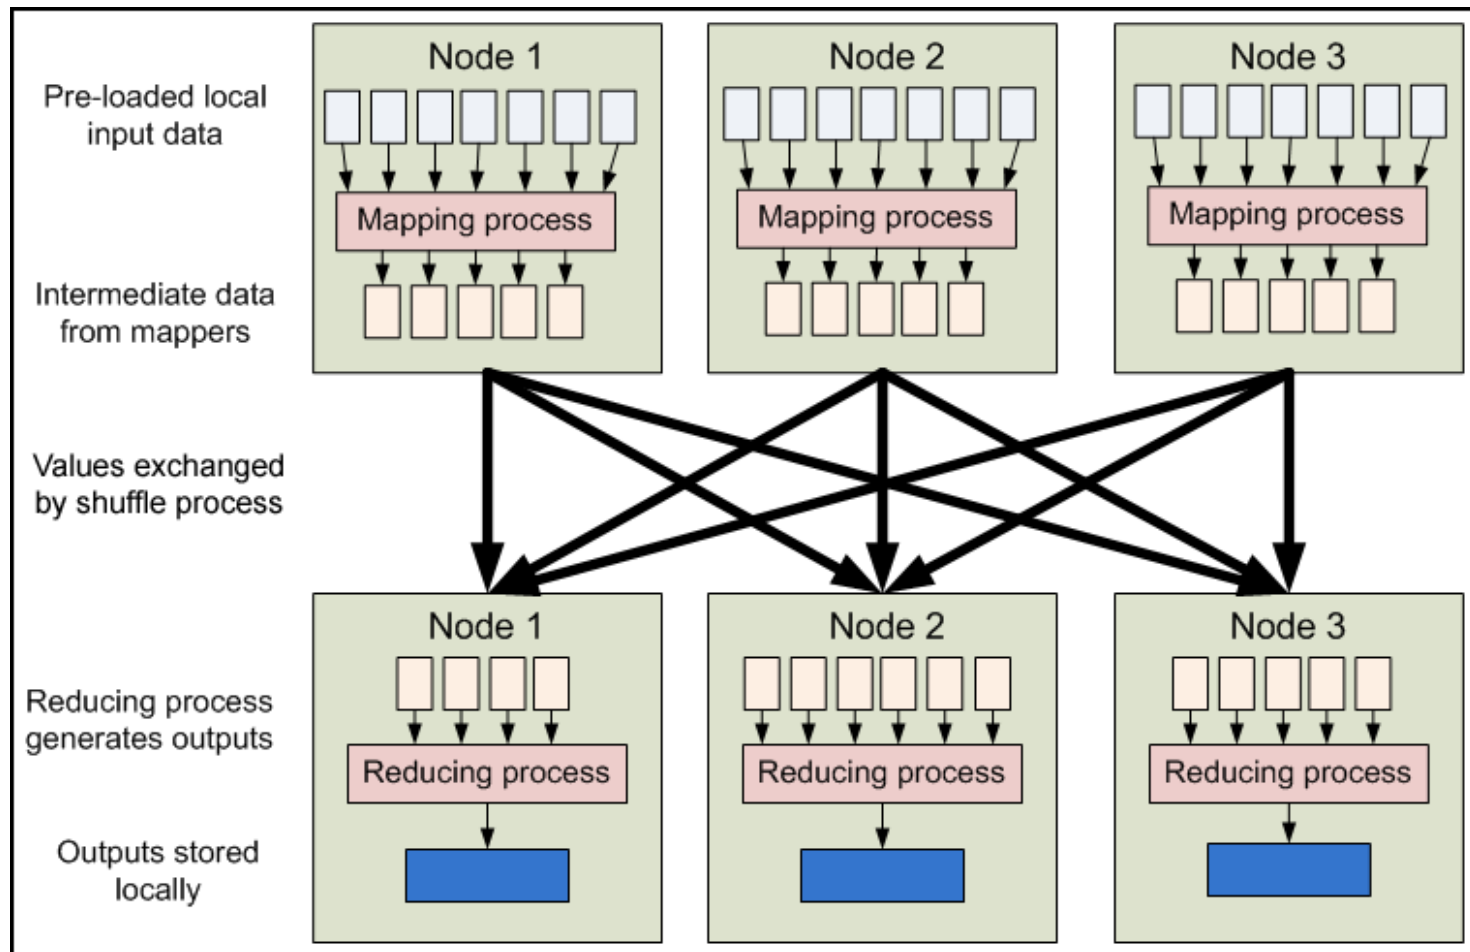
How is reduce() in Python evaluated?

- $z = \text{reduce}(f, \text{list})$ where f is add function
- Initially, z (an accumulator) is set to $\text{list}[0]$
- Next, repeat $z = \text{add}(z, \text{list}[i])$ for each $i > 0$
- Return final z
- Example: $z = \text{reduce}(\text{add}, [1, 2, 3])$
 - $i = 0, z = 1; i = 1, z = 3; i = 2, z = 6$

Hadoop MapReduce

- Map
 - $\langle k, v \rangle \Rightarrow$ list of $\langle k', v' \rangle$
- Reduce:
 - $\langle k', \text{list of } v' \rangle \Rightarrow$ list of $\langle k'', v'' \rangle$
- Write MapReduce programs on Hadoop
 - Using Java

MapReduce



WordCount: mapper

Object can be replaced with LongWritable

```
public class WordCount {  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(Object key, Text value, Context context  
                        ) throws IOException, InterruptedException {  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
}
```

Data types of input key-value

Data types of output key-value

Key-value pairs with specified data types

The diagram illustrates the data types used in the WordCount mapper. Red boxes highlight the generic type parameters in the `Mapper` class declaration and the `map` method signature. Blue arrows point from descriptive text to these boxes and to the `context.write` statement.

- Data types of input key-value:** Points to the `Object` and `Text` parameters in the `Mapper` class declaration.
- Data types of output key-value:** Points to the `Text` and `IntWritable` parameters in the `Mapper` class declaration.
- Key-value pairs with specified data types:** Points to the `context.write(word, one);` statement.

WordCount: reducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

Data types of input key-value

Data types of output key-value

A list of values

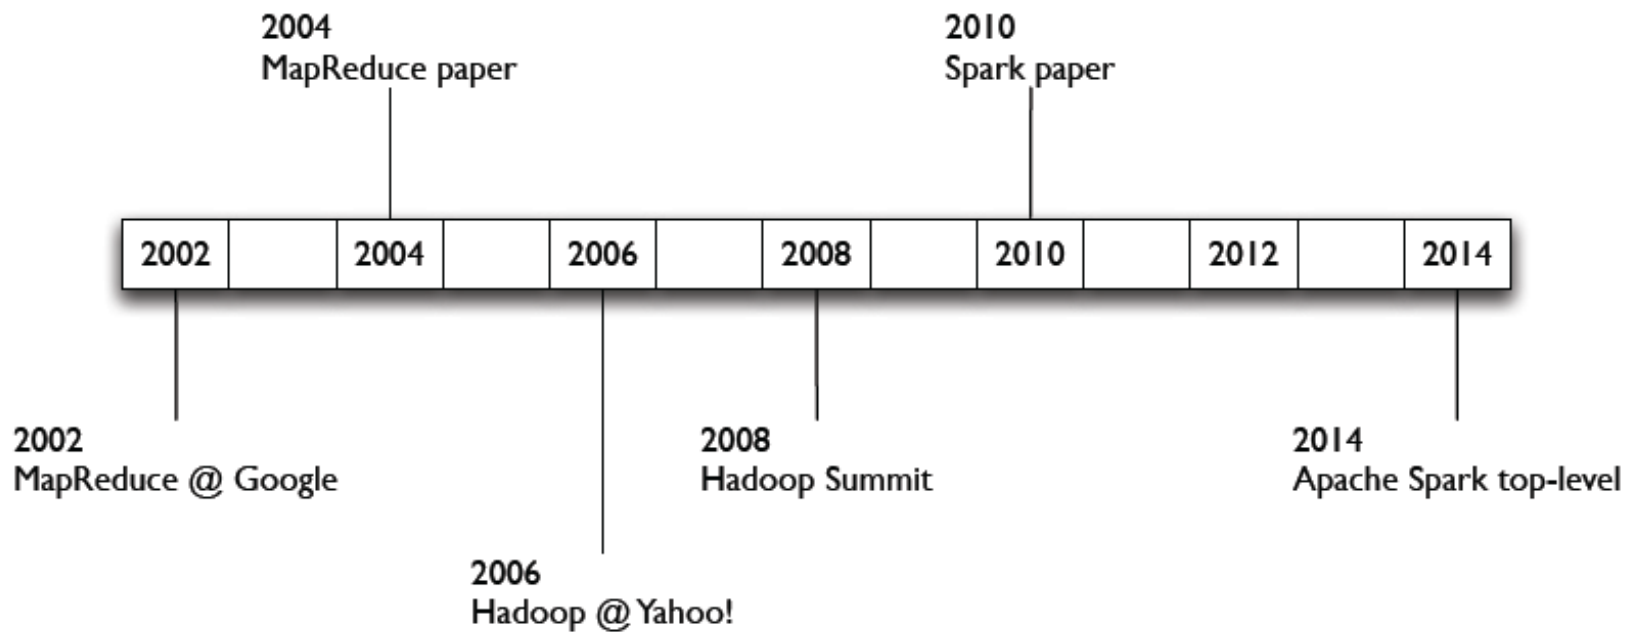
Limitations of MapReduce

- Acyclic dataflow
 - difficult to reuse computed data
 - ill-suited for iterative algorithms
 - e.g., k-means, PageRank, logistic regression
- Batch processing
 - Not suitable for interactive data analysis

In-memory MapReduce (Spark)

- Key concepts
 - RDD (resilient distributed dataset)
 - Transformations
 - Actions

Apache Spark: history




Spark

- Support working sets through RDD
 - reuse & fault-tolerance
- 10x faster than Hadoop in iterative jobs
- Interactively explore 39GB with sub-second response time

Spark

- Combine SQL, streaming, and complex analytics



Spark

- Run on Hadoop, Cassandra, HBase, etc.



wc.py

```
from pyspark import SparkContext
from operator import add

sc = SparkContext(appName="inf551")

lines = sc.textFile('hello.txt')

counts = lines.flatMap(lambda x: x.split(' ')) \
               .map(lambda x: (x, 1)) \
               .reduceByKey(add)

output = counts.collect()

for v in output:
    print '%s, %s' % (v[0], v[1])
```

Major topics

- Storage systems
- File systems & file formats
- Database management systems
- Big data solution stack
- Data warehousing (if time permits)



Data warehousing

- Multidimensional data model
 - Star vs snowflake schema
- OLAP operations: rollup, drill-down, etc.
- Materialized views
- Index (we will cover this if time permits)
 - Bitmap
 - Run-length encoding
 - Join index

What is a Warehouse?


- Collection of diverse data
 - subject oriented, e.g., sales
 - aimed at executive, decision maker
 - often a copy of operational data
 - with value-added data (e.g., summaries, history)
 - integrated
 - time-varying: historical data, discovering trend
 - non-volatile: once in warehouse, data do not change



What is a Warehouse?





- Collection of tools
 - gathering data
 - cleansing, integrating, ...
 - querying, reporting, analysis
 - data mining
 - monitoring, administering warehouse

Warehouse Architecture




Why a Warehouse?

- Two Approaches to Integration:
 - Warehouse (Eager)
 - Query-Driven (Lazy)



Query-Driven Approach



Advantages of Warehousing

- High query performance
- Local processing at sources unaffected
- Can operate when sources unavailable
- Extra information at warehouse
 - Modify, summarize (store aggregates)
 - Add historical information

Disadvantages of Warehousing

- May require hefty storage space
- Need to decide what to store in advance
- Can only query data stored in warehouse
 - Data get stale
- Must detect source changes & update warehouse

Advantages of Query-Driven

- No need to copy data, less storage
- No need to purchase data
- More up-to-date data
- Query needs can be unknown

Disadvantages of Query-Driven

- Inefficient/delay in query processing
 - source unreliable
 - slow network
 - expensive translation, filtering, merging
- Sources might not permit ad-hoc queries
 - Examples?

OLTP vs. OLAP

- OLTP: On Line Transaction Processing
 - Describes processing at operational sites (order entry in POS/online, banking transactions, etc.)
- OLAP: On Line Analytical Processing
 - Describes processing (answering analytical queries: aggregation, rollup/drilldown, slice/dice, etc.) at warehouse

OLTP vs. OLAP

OLTP

- Mostly updates
- Many small transactions
- Mb-Tb of data
- Raw data
- Clerical users
- Up-to-date data
- Consistency, recoverability critical

OLAP

- Mostly reads
- Queries long, complex
- Gb-Tb of data
- Summarized, consolidated data
- Decision-makers, analysts as users
- Historical data
- Query performance critical


Big data ETL & Warehousing

- Apache pig
 - Focus on ETL & data transformations
 - Compile transformations into MapReduce jobs
 - Pig latin script is **procedural** (step-by-step)
- Apache Hive
 - **Declarative** HiveQL (SQL-like)
 - Queries are turned into MapReduce jobs


After class: sign up for AWS




Dashboard after logging on

 **AWS** ▾ **Services** ▾ **Edit** ▾ Wensheng Wu ▾ N. Virginia ▾


Quick Starts HIDE




Build a web app
[Start now](#)




Launch a Virtual Machine (EC2 Instance)




Back up your files
[Learn more](#)



Build a back end for your mobile app




Host a static website




Analyze big data
[Learn more](#)


Shortcuts and Recently Viewed Services




DynamoDB



EC2



IAM



Cognito

GETTING STARTED

Read our [documentations](#) or view our [training](#) to learn about AWS.

AWS CONSOLE MOBILE APP

View your resources on the go with our AWS Console mobile app, available from [Amazon Appstore](#), [Google Play](#), and [iTunes](#).

AWS MARKETPLACE

[Find and buy software](#) with 1-Click, and pay by the hour.

Lab session

- Task: Setting up an EC2 instance
- Details: see lab session slides to be posted...