

File Systems

INF 551

Wensheng Wu

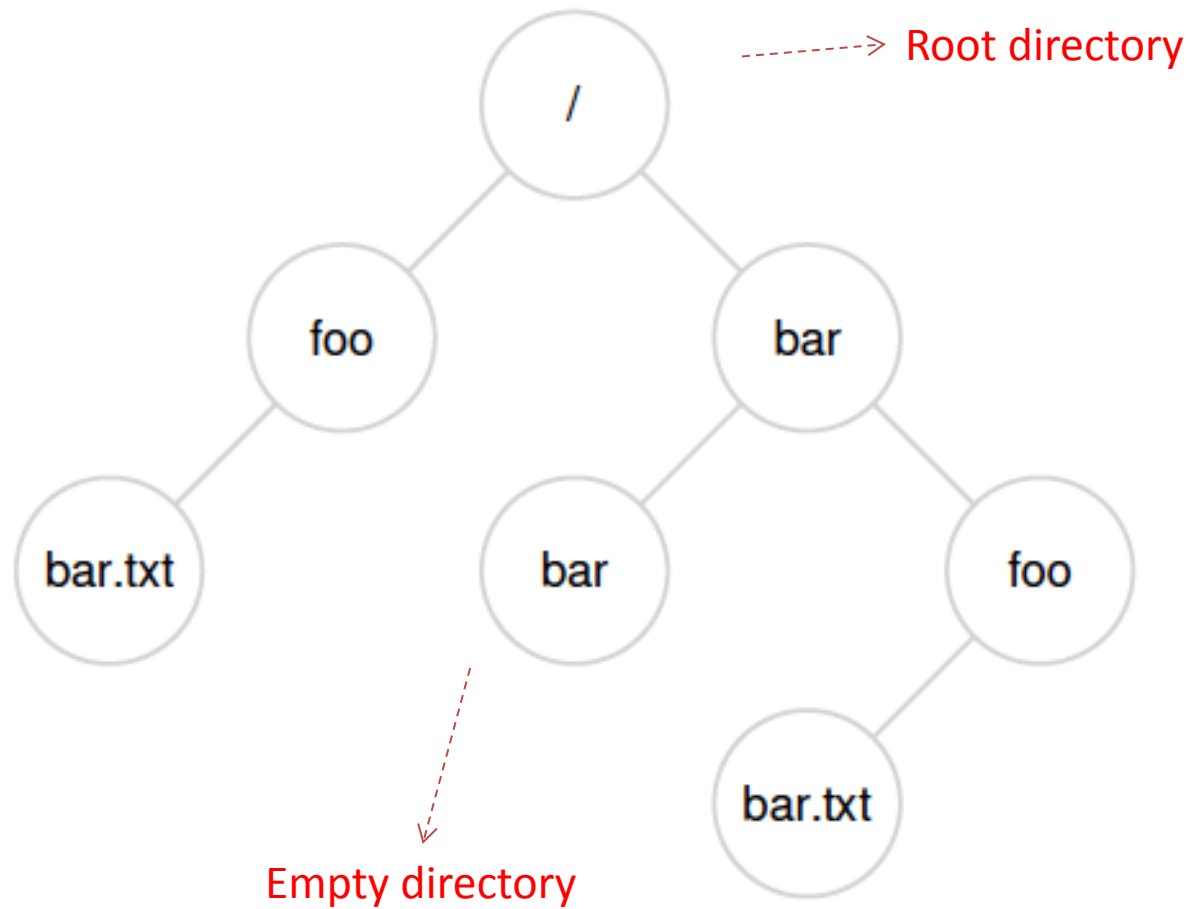
Roadmap

- Files and directories
 - CRUD operations
- How to implement them
 - Data structures
 - Access methods

Files and directories

- File: stored in blocks on storage device
 - Has user defined name: hello.txt
 - & low-level name, e.g., inode number: 410689
- Files are organized into directories (folders)
 - each may have a list of files and/or subdirectories
 - That is, directories can be nested


Example



Operations on files

- Create
- Read
- Update
- Delete

Create

- User interface, e.g., via GUI
 - Implementation, e.g., via a C program with a call to system function `open()`
 - `int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC);`
 - Open with flags indicating the specifics
 - `O_CREATE`: create a file
 - `O_WRONLY`: write only
 - `O_TRUNC`: remove existing contents if exists
-  `|`: Bitwise OR operator

File descriptor

- Note `open()` returns a file descriptor
 - Typically an integer
 - Reserved fds: `stdin` 0, `stdout`, 1, `stderr` 2

Read

- `read(fd, buffer, size)`
 - Read from file "fd" *<size>* number of bytes
 - And store them in *buffer*
- Read starts from the current offset of fd
 - Initially 0

Write

- `write(fd, buffer, size)`
 - Write to file *fd* *<size>* number of bytes stored in buffer
 - Also start writing from the current offset

Random read and write

- `off_t lseek(int fd, off_t offset, int whence)`
 - If `whence` is `SEEK_SET`, the offset is set to `<offset>` bytes from the beginning of file
 - If `whence` is `SEEK_CUR`, the offset is set to its current location plus `<offset>` bytes
 - If `whence` is `SEEK_END`, the offset is set to the size of the file plus `<offset>` bytes (typically offset is negative, e.g., -8 for 8 bytes from the end)
- `whence`: from where

```
#define BUF_SIZE 8192
```

```
int main(int argc, char* argv[]) {
```

```
    int input_fd, output_fd;    /* Input and output file descriptors */
    ssize_t ret_in, ret_out;    /* Number of bytes returned by read() and write() */
    char buffer[BUF_SIZE];     /* Character buffer */
```

```
    /* Are src and dest file name arguments missing */
```

```
    if(argc != 3){
        printf ("Usage: cp file1 file2\n");
        return 1;
    }
```

→ Copy a file

```
    /* Create input file descriptor */
```

```
    input_fd = open (argv [1], O_RDONLY);
    if (input_fd == -1) {
        perror ("open");
        return 2;
    }
```

"0" starts an octal number
=> permissions:

110 (owner) rw-
100 (group) r--
100 (others) r--

```
    /* Create output file descriptor */
```

```
    /* WRONLY will truncate file to zero length if exists */
```

```
    output_fd = open(argv[2], O_WRONLY | O_CREAT, 0644);
    if(output_fd == -1){
        perror("open");
        return 3;
    }
```

→

```
    /* Copy process */
```

```
    while((ret_in = read (input_fd, &buffer, BUF_SIZE)) > 0){
        ret_out = write (output_fd, &buffer, (ssize_t) ret_in);
        if(ret_out != ret_in){
            /* Write error */
            perror("write");
            return 4;
        }
    }
```

→ Pointer to a character array

File permission mode

```
Vincent@Vincent-PC ~/usc/551-fa16/Amazon
$ copy2
Usage: cp file1 file2

Vincent@Vincent-PC ~/usc/551-fa16/Amazon
$ copy2 copy2.c copy2-a.c

Vincent@Vincent-PC ~/usc/551-fa16/Amazon
$ ls -l
total 95
-rwxrwx---+ 1 Vincent None      3 Aug 30 18:04 a.txt
-rwxrwx---+ 1 Vincent None      0 Aug 30 18:04 a.txt~
-rwxrwx---+ 1 Vincent None 1568 Jan 31 2016 copy2.c
-rwxrwxr-x+ 1 Vincent None 64289 Sep 10 11:45 copy2.exe
-rw-r--r---+ 1 Vincent None 1568 Sep 10 11:45 copy2-a.c
-rwxrwx---+ 1 Vincent None  426 Aug 31 15:18 HelloWorld.class
-rwxrwx---+ 1 Vincent None  239 Aug 30 18:02 HelloWorld.java
-r-----+ 1 Vincent None 1698 Aug 23 17:18 inf551.pem
-rwxrwx---+ 1 Vincent None 1464 Aug 23 20:56 inf551.ppk
-r-----+ 1 Vincent None 1694 Aug 31 14:48 inf551-a.pem
-r-----+ 1 Vincent None 1698 Aug 31 17:28 inf551-b.pem
-rwxrwx---+ 1 Vincent None 1464 Aug 31 17:39 inf551-b.ppk
```

rw-r--r-

=> 110 (owner permission) 100 (group) 100 (others)

Resources for system calls

- https://en.wikipedia.org/wiki/System_call
- open:
[https://en.wikipedia.org/wiki/Open \(system call\)](https://en.wikipedia.org/wiki/Open_(system_call))
- read:
[https://en.wikipedia.org/wiki/Read \(system call\)](https://en.wikipedia.org/wiki/Read_(system_call))
- write:
[https://en.wikipedia.org/wiki/Write \(system call\)](https://en.wikipedia.org/wiki/Write_(system_call))
- close:
[https://en.wikipedia.org/wiki/Close \(system call\)](https://en.wikipedia.org/wiki/Close_(system_call))

Resources for system calls

- `man -S 2 read`
 - Find it in the Section 2 of the manual

Install gcc on EC2

- `sudo yum groupinstall "Development Tools"`
 - Will install other dev. tools too
 - E.g., perl, bison, flex, automake, autoconf
- Usage:
 - `gcc -o copy2 copy2.c`

File and directory

- When creating a file
 - Bookkeeping data structure (inode) created: recording size of file, location of its blocks, etc.
 - Linking a human-readable name to the file
 - Putting the link in a directory

Info about file (stored in inode)

```
struct stat {  
    dev_t st_dev; /* ID of device containing file */  
    ino_t st_ino; /* inode number */  
    mode_t st_mode; /* protection */  
    nlink_t st_nlink; /* number of (hard) links */  
    uid_t st_uid; /* user ID of owner */  
    gid_t st_gid; /* group ID of owner */  
    dev_t st_rdev; /* device ID (if special file) */  
    off_t st_size; /* total size, in bytes */  
    blksize_t st_blksize; /* blocksize for filesystem I/O */  
    blkcnt_t st_blocks; /* number of blocks allocated */  
    time_t st_atime; /* last time file content was examined */  
    time_t st_mtime; /* last time file content was changed */  
    time_t st_ctime; /* last time inode was changed */  
};
```

inode

- Stores metadata/attributes about the file
- Also stores locations of blocks holding the content of the file

Example

- a.txt

abc def

abc def

abc def

Device id

Block size

of blocks allocated

Inode #

of (hard) links

```
[ec2-user@ip-172-31-52-194 inf551]$ stat a.txt
  File: 'a.txt'
  Size: 24          Blocks: 8          IO Block: 4096   regular file
Device: ca01h/51713d Inode: 410837       Links: 1
Access: (0770/-rwxrwx---)  Uid: ( 500/ec2-user)   Gid: ( 500/ec2-user)
Access: 2016-09-10 23:57:57.757982711 +0000
Modify: 2016-09-10 23:57:57.869981750 +0000
Change: 2016-09-10 23:57:57.869981750 +0000
 Birth: -
[ec2-user@ip-172-31-52-194 inf551]$
```

Access permission

User id

Group id

Working with directories

- Create: `mkdir()` system call
 - Used to implement command, e.g., `mkdir xyz`
- Read: `opendir()`, `readdir()`, `closedir()`
 - `ls xyz`
- Delete: `rmdir()`

Roadmap

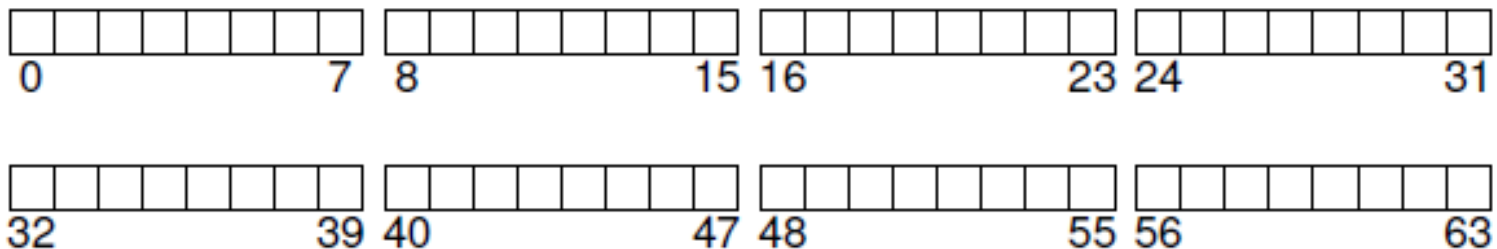
- Files and directories
 - CRUD operations
- Implementation
 - Data structures: how to organize the blocks
 - Access methods: map system calls to operations on data structures

Organization of blocks

- Array-based
 - Disk consists of a list of blocks
 - We will assume this
- Tree-based, e.g., SGI XFS
 - Blocks are organized into **variable-length** extents
 - Use B+-tree to quickly find free extents

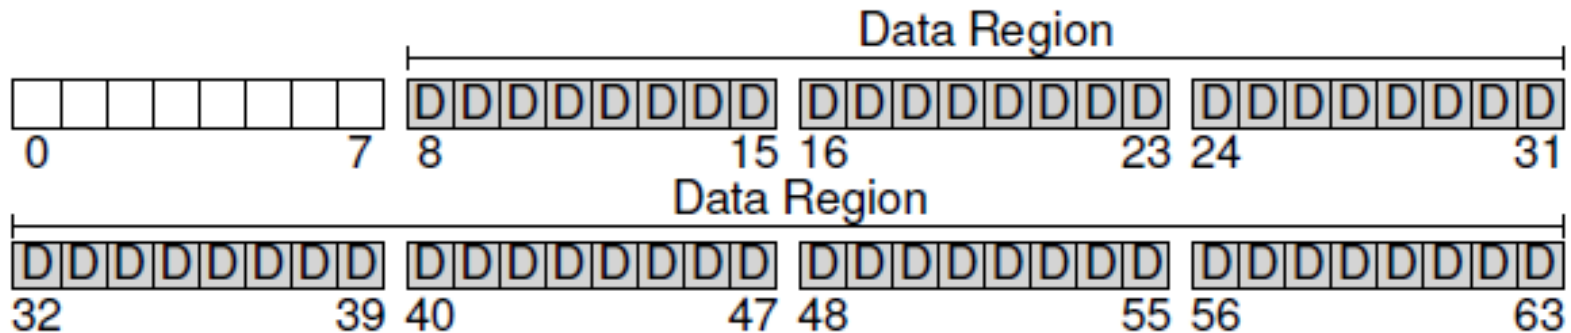
Blocks

- Consider a disk with 64 blocks
 - 4KB/block
 - 512B/sector (we assume this in this lecture)
- So there are $2^{12}/2^9 = 2^3 = 8$ sectors/block
 - Capacity of disk = $64 * 4KB = 256KB$



Data region

- 56 blocks used to store data (files)
 - Blocks # 8 – 63

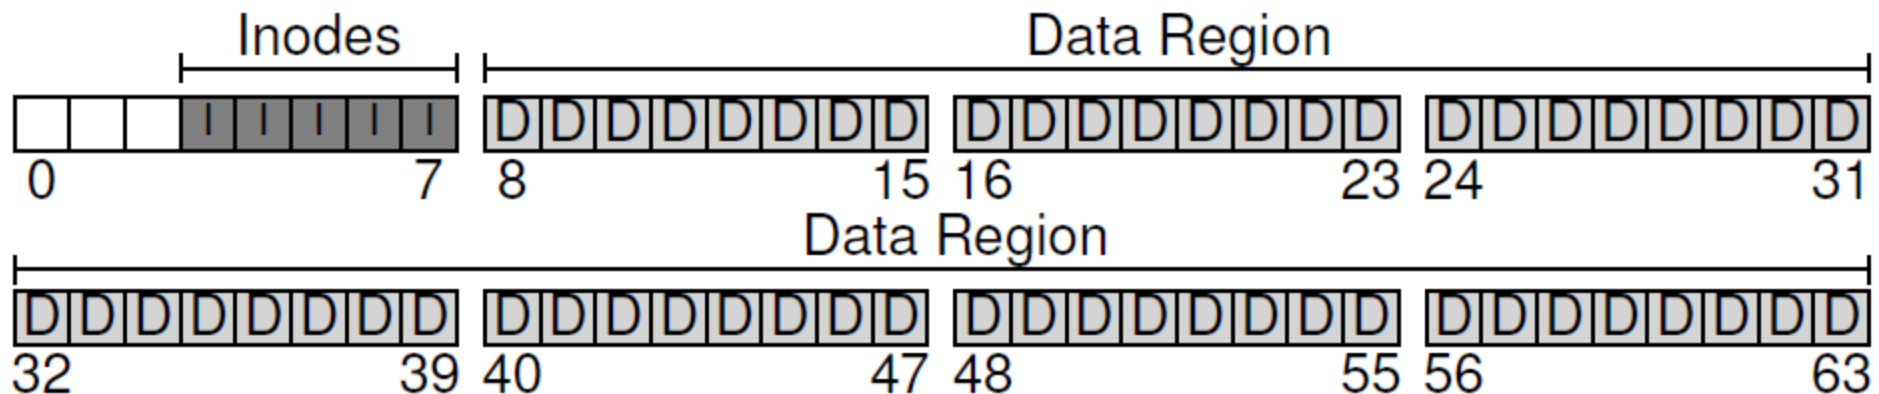


Metadata

- For each file, file system keeps track of
 - Location of the blocks that comprise the file
 - Size of the file
 - Owner and access rights
 - Access and modify times
 - Etc. (see the stat struct a couple of slides back...)
- These metadata are stored in an **inode** (index node)

inodes

- Index nodes
- Stored in blocks #3 -- #7 (i.e., 5 blocks)
- Together they are called inode table



How many inodes are there?

- 256 bytes/inode
- 5 blocks, 4KB/block

=> 16 inodes/block ($4K/256 = 2^{12}/2^8$)

=> 5 blocks, $5 * 16 = 80$ inodes

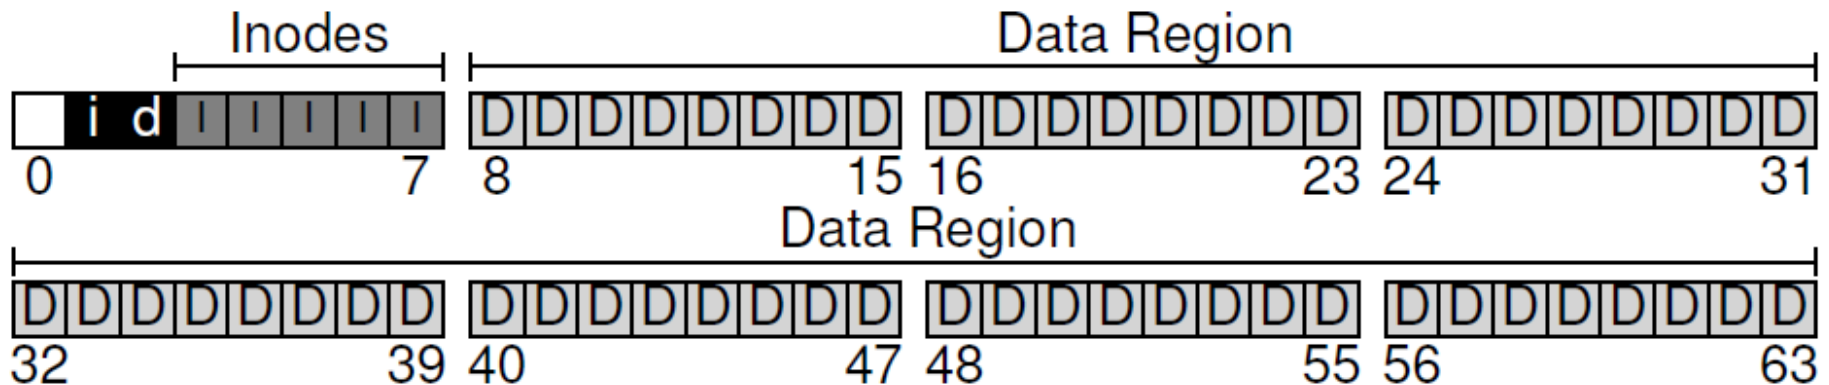
=> File system can store at most 80 files

Free space management using bitmaps

- Bitmap: a vector of bits
 - 0 for free (inode/block), 1 for in-use
- Inode bitmap (imap)
 - keep track of which inodes in the inode table are available
- Data bitmap (dmap)
 - Keep track of which blocks in data region are available

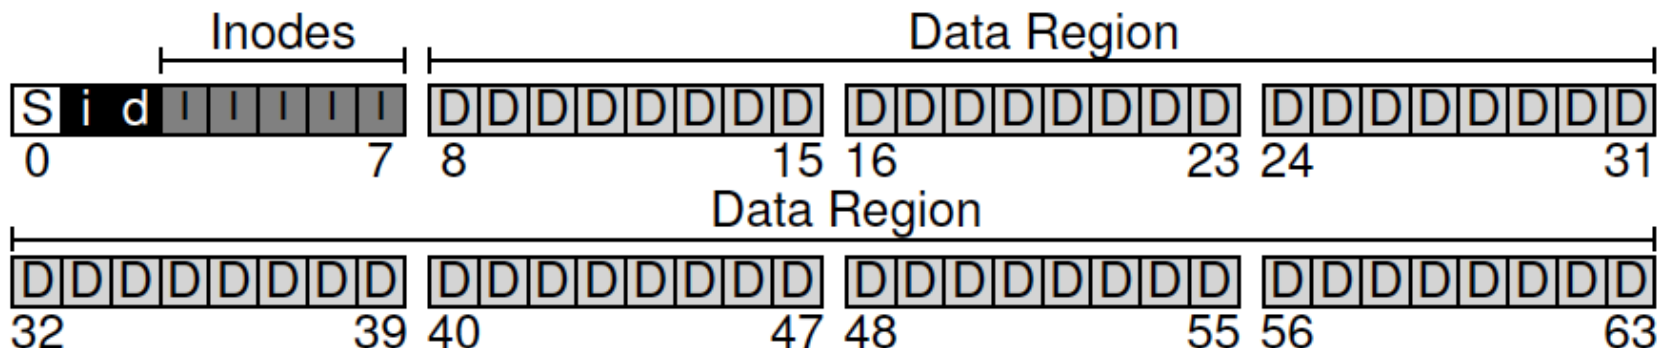
Bitmaps

- Each bitmap is stored in a block
 - Block "i": keep track of 80 inodes (could track 32K)
 - Block "d": keep track of the 56 data blocks



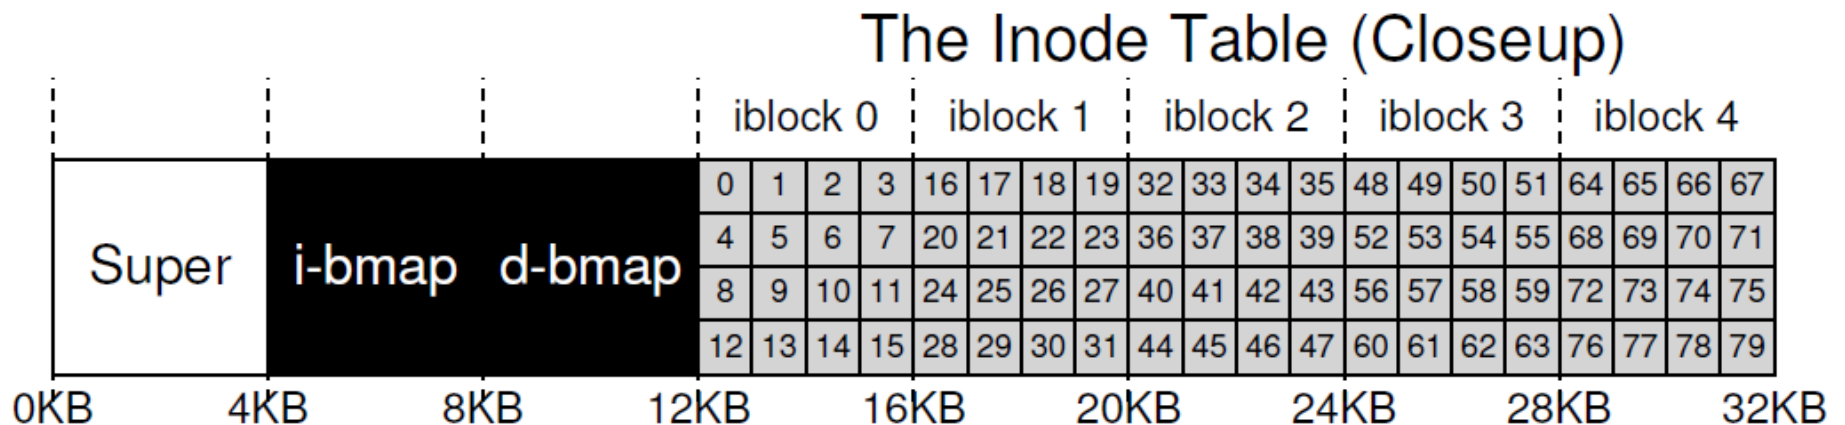
Superblock

- Track where i/d blocks and inode table are
 - E.g., inode table starts at block 3; there are 80 inodes and 56 data blocks, etc.
- Indicate type of file system
- Will be read first when file system is mounted



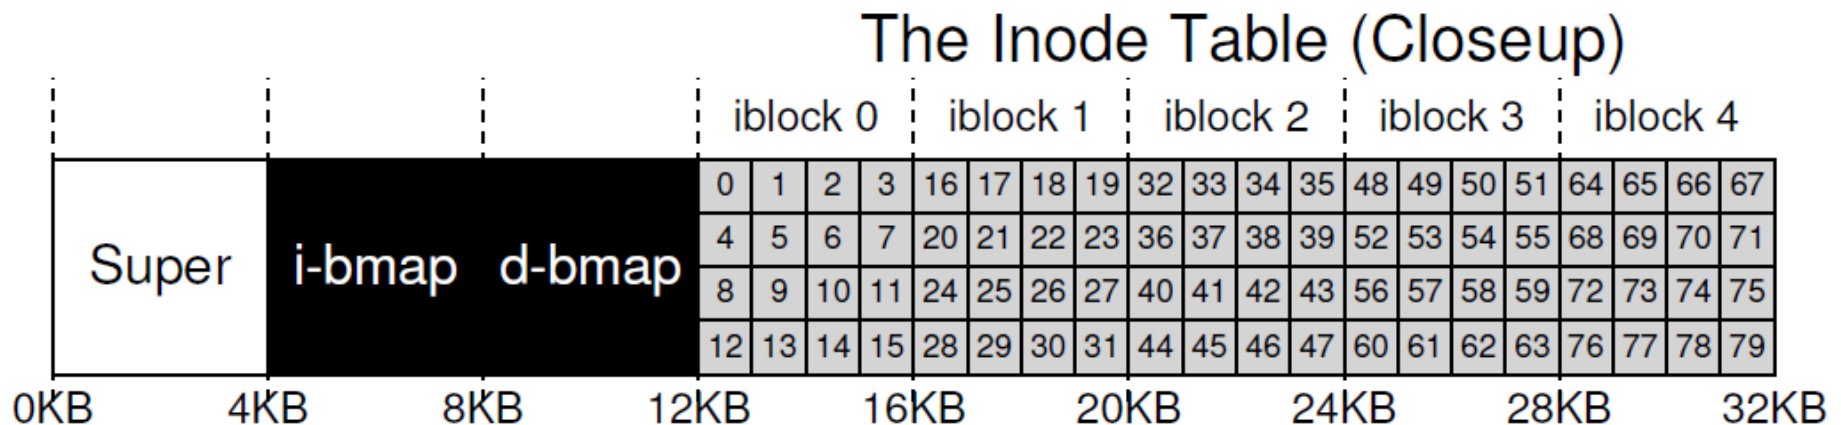
inumber

- Each inode is identified by a number
 - Low-level number of file name
- Can figure out location of inode from inumber



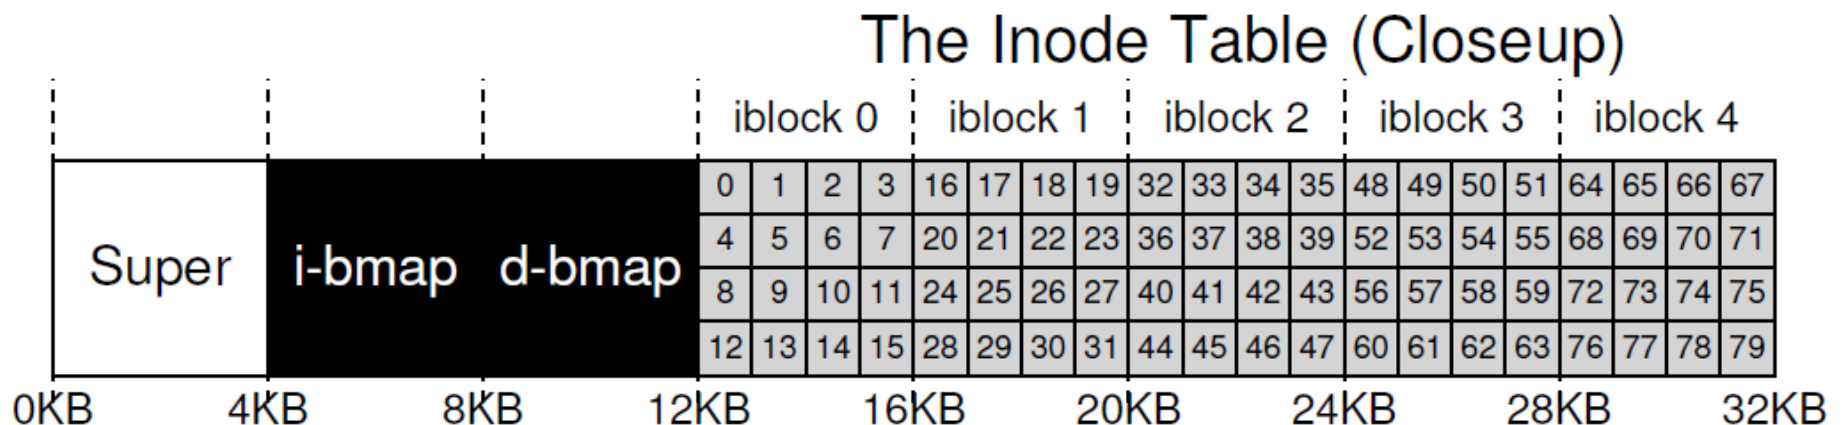
inumber => location

- inumber = 32
=> address: offset in bytes from the beginning
=> which sector?



inumber => location of inode

- Address: $12K + 32 * 256 = 20K$
- Sector #: $20K / 512 = 40$
 - more generally
 - $\lfloor (\text{inodeStartAddress} + \text{inumber} * \text{inode size}) / \text{sector size} \rfloor$



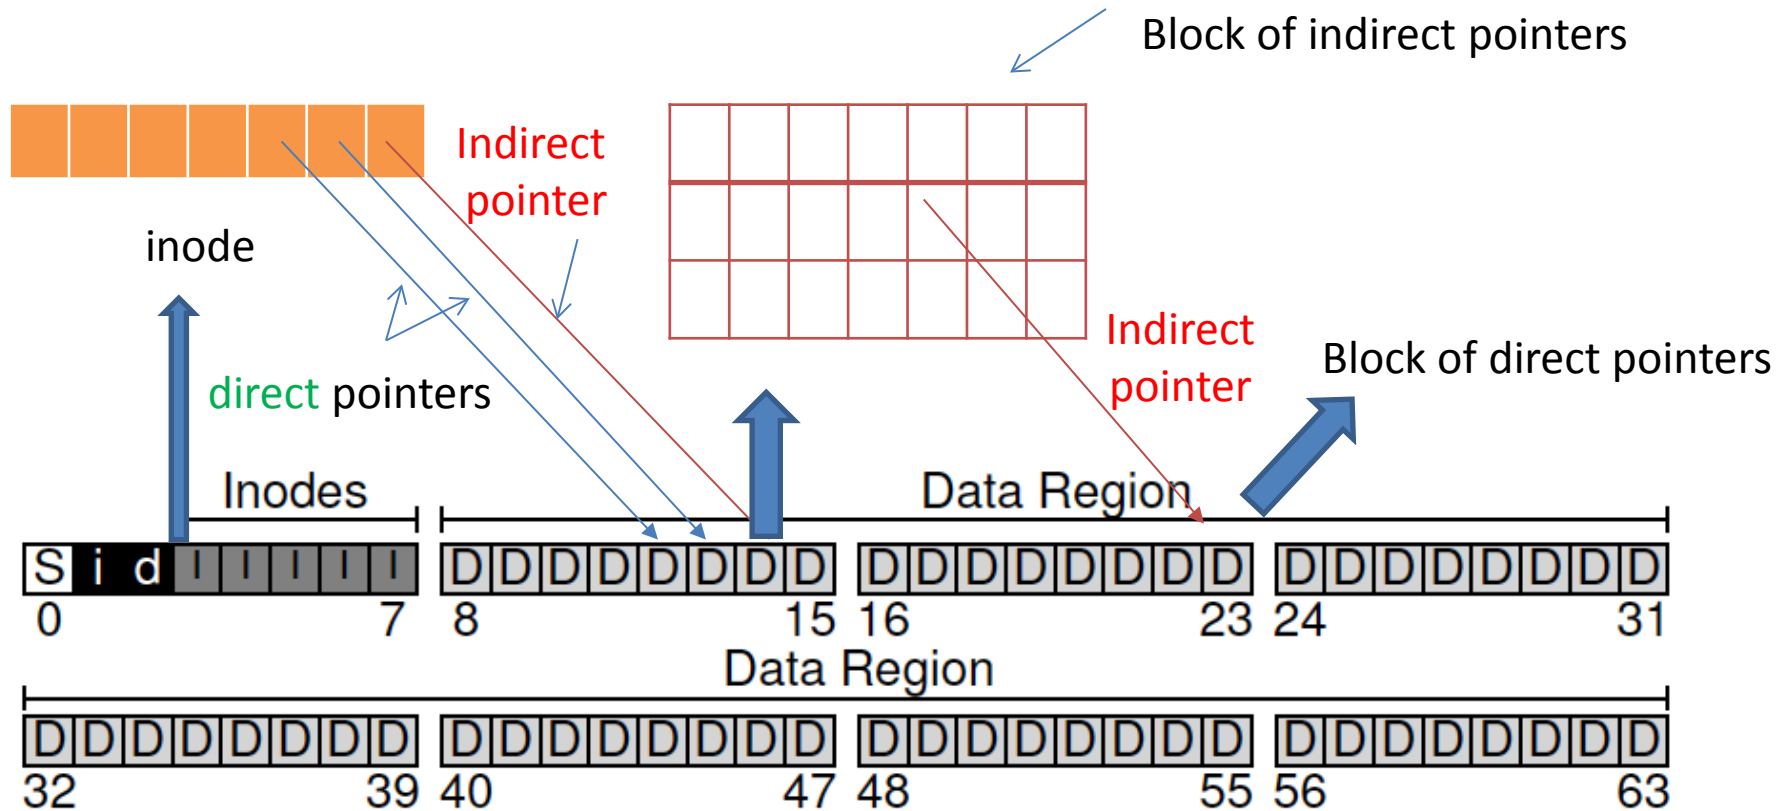
inode => location of data blocks

- A number of direct pointers
 - E.g., 8 pointers, each points to a data block
 - Enough for $8 * 4K = 32K$ size of file
- Also has a slot for indirect pointer
 - Pointing to a data block storing direct pointers
 - Assume 4 bytes for block address (e.g., represented in CHS), so 1024 pointers/block
 - Now file can have $(8 + 1024)$ blocks or 4,128KB

Multi-level index

- Pointers may be organized into multiple levels
 - Indirect pointer (as in previous slide)
 - Inode (pointer1, pointer2, ..., indirect pointer)
 - **Indirect** pointer -> a block of **direct** pointers
 - Double indirect pointers
 - Inode (pointer1, pointer2, ..., indirect pointer)
 - **Indirect** pointer -> a block of **indirect** pointers instead
 - > each points to a block of **direct** pointers
 - Triple indirect pointers
 - **Indirect** pointer -> a block of **indirect** pointers
 - > each points to a block of **indirect** pointers
 - > each points to a block of **direct** pointers

Double Indirect Pointers



Advantages of multi-level index

- Grow to more levels as needed
- Direct pointers handle most of the cases
 - Many files are small

Directory organization

- Directory itself stored as a file
- For each file in the directory, it stores:
 - name, inumber, record length, string length

<code>inum</code>	<code>reclen</code>	<code>strlen</code>	<code>name</code>
5	4	2	.
2	4	3	..
12	4	4	foo
13	4	4	bar
24	8	7	foobar

Actual length



Record length vs string length

- String length = # of characters in file name + 1
(for \0: end of string)
- Record length \geq string length
 - Due to entry reuse

<code>inum</code>	<code>reclen</code>	<code>strlen</code>	<code>name</code>
5	4	2	.
2	4	3	..
12	4	4	foo
13	4	4	bar
24	8	7	foobar


Reusing directory entries

- If file is deleted (using `rm` command) or a name is unlinked (using `unlink` command)
 - File is finally deleted when its last (hard) link is removed
- Then `inumber` in its directory entry set to 0 (reserved for empty entry)
 - So we know it can be reused

Storing a directory

- Also as a file with its own inode + data block
- inode:
 - file type: directory (instead of regular file)
 - pointer to block(s) in data region storing directory entries

Roadmap

- Files and directories
 - CRUD operations
 - Implementation
 - Data structures: how to organize blocks, e.g., into array/tree
 - **Access methods**: turn system calls to operations on data structures
- 

Open for read

- `fd = open("/foo/bar", O_RDONLY)`

↑
mode

bar Inode



foo



foo



root

content

Inode

content →

root



mount

Inode

(super block)

Open for read

- `fd = open("/foo/bar", O_RDONLY)`
 - Need to locate inode of the file `"/foo/bar"`
 - Assume inumber of root, say 2, is known (e.g., when the file system is mounted)

Open for read

1. Read inode and content of / (2 reads)
 - Look for "foo" in / -> foo's inumber
2. Read inode and content of /foo (2 reads)
 - Look for "bar" in /foo -> bar's inumber
3. Read inode of /foo/bar (1 read)
 - Permission check + allocate file descriptor

Cost of open()

- Need 5 reads of inode/data block

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[1]
open(bar)			read	read	read	read	read			
read()					read			read		
read()					write				read	
read()					read					
read()					write					
					read					
					write					read

Reading the file

- `read(fd, buffer, size)`
 - Note `fd` is maintained in per-process open-file table
 - Table translates `fd` -> inumber of file

File-open table per process

File descriptor	File name	Inumber	Position offset	...
3	/foo/bar	32382	0	
4	/foo/more	48482	512	...

Reading the file

- `read(fd, buffer, size)`
 1. Consult bar's inode to locate its 1st block
 2. Read the block
 3. Update inode with newest file access time
 4. Update open-file table with new offset
 5. Continue steps 2, 3, 4 until done
 6. Deallocate file descriptor

Cost for reading a block

- 3 I/O's:
 - read inode, read data block, write inode

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[1]	2
open(bar)			read	read	read	read	read				
read()					read			read			
read()					write						
read()					read				read		
read()					read						
read()					write					read	

Open for write

- `int fd = open("/foo/bar", O_WRONLY)`
 - Assume bar is a **new** file under foo
 - (note the difference from reading chapter!)

Open for write

- `int fd = open("/foo/bar", O_WRONLY)`
 1. Read '/' inode & content
 - obtain foo's inumber
 2. Read '/foo' inode & content
 - check if bar exists

Open for write

3. Read imap, to find a free inode for bar
4. Update imap, setting 1 for allocated inode
5. Write bar's inode

Open for write

6. Update foo's content block

- Adding an entry for bar

7. Update foo's inode

- Update its modification time

Cost for "open for write"

- `int fd = open("/foo/bar", O_WRONLY)`
- Need 9 I/O's

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create()			read							
						read				
				read						
							read			
		read								
		write								
					write					
							write			
				write						

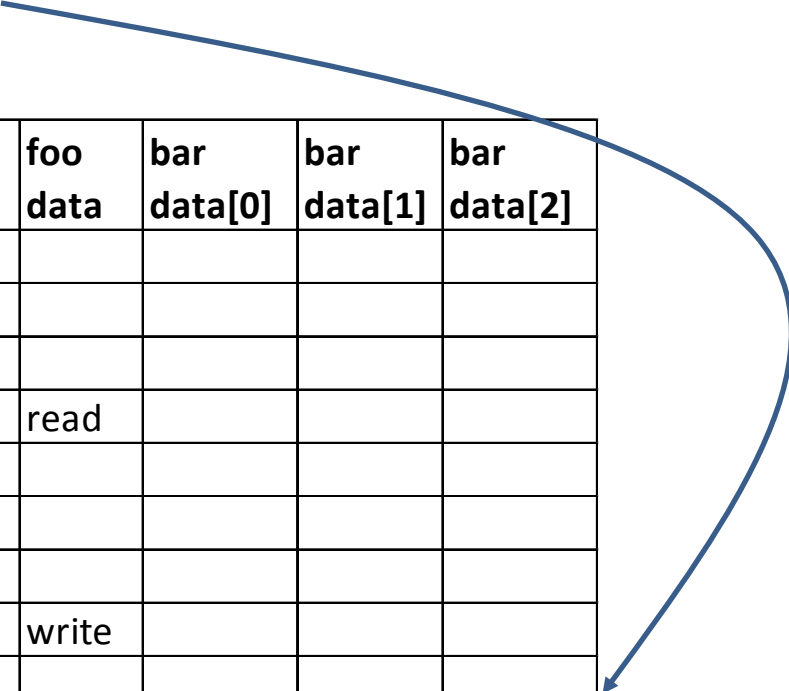
Writing the file: /foo/bar

1. Read inode of bar (by looking up its inumber in the file-open table)
2. Allocate new data block
 - Read and write bmap
3. Write to data block of bar
4. Update bar inode
 - new modified time, add pointer to block

Cost of writing /foo/bar

- 5 I/O's for write a block

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create()			read							
						read				
				read						
							read			
		read								
		write								
					write					
				write			write			
write()					read					
	read									
	write									
								write		
					write					



Caching for read

- First read may be slow
 - But subsequent ones will speed up
- Good idea to cache popular blocks
 - e.g., determined via LRU strategy

Buffering for delayed write

- Improve write performance via:
 - Batching (e.g., two updates to the same imap)
 - Scheduling (reordering for better performance)
 - Avoiding writes (if file created, then quickly deleted)
- Problem: update may be lost when system crashes

Example file systems

- NTFS
 - New technology file system, Microsoft proprietary
- FAT
 - File allocation table
 - FAT 16, 32, ...
 - 32 bits = # of sectors a file can occupy
 - 512B/sector => 2TB limit on file size
 - 4KB/sector => 16TB limit
- Ext4
 - fourth extended file system, common in Linux