

## INF 551 – Fall 2017 (Afternoon section)

## Quiz 11: Query execution (10 points), 15 minutes

**Note: in all questions, we ignore the cost of outputting the result of join.**

1. [5 points] Suppose  $M = 101$  blocks,  $B(R) = 2,000$  blocks,  $B(S) = 6,000$  blocks. Suppose we use 100 blocks in sorting. Describe the steps (what each step does, what does it output, what is the size of the output) of computing  $R \bowtie S$  using the **simple sort-merge** join algorithm. What is the cost of the algorithm?

Step 1: Sort R: load 100 blocks of R at a time, sort them and write back on disk. This generates 20 runs, each of size 100 blocks.

Cost =  $2B(R) = 4000$

Step 2: Load 20 runs of R at a time, and merge them. It generates 1 run and the run size is 2000 pages

Cost =  $2B(R) = 4000$

Step 3: Sort S: load 100 blocks of S at a time, sort them and write back on disk. This generates 60 runs, each of size 100 blocks

Cost =  $2B(S) = 12000$

Step 4: Load 60 runs of S at a time, and sort them. It generates 1 run and run size is 6000 pages

Cost =  $2B(S) = 12000$

Step 5: Joining R and S using two-way merging, since only one run from R and one from S.

Cost =  $B(R) + B(S) = 8000$

Total cost = 40000.

2. [5 points] Suppose  $M = 101$  blocks (i.e., pages),  $B(R) = 20,000$  blocks,  $B(S) = 50,000$  blocks. Describe the steps of computing  $R \bowtie S$  using the **partitioned-hash** join algorithm. What is the cost of the algorithm?

Step 1: Hash R into 100 buckets using hash function  $h$ , and send all buckets to disk. 200 blocks/bucket.

Step 2: Hash S into 100 buckets using hash function  $h$ , and send all buckets to disk. 500 blocks/bucket.

The memory is not enough to hold an entire bucket of R (200 blocks) or an entire bucket of S (500 blocks).

Step 3: Hash each bucket of R into 100 buckets using hash function  $h'$ , and send all buckets to disk. 2 blocks/bucket.

Step 4: Hash each bucket of S into 100 buckets using hash function  $h'$ , and send all buckets to disk. 5 blocks/bucket.

Step 5: Join corresponding buckets from R and S (i.e., with the same bucket number after both hash).

Total cost:  $5B(R) + 5B(S) = 350,000$