

INF 551 – Fall 2018 (Afternoon)

Quiz 7: Query execution (10 points)

10 minutes

Consider natural-joining two relations $R(A, B)$ and $S(B, C)$ using the **partitioned hash join** algorithm. Suppose $M = 101$ pages, $B(R) = 5,000$ blocks, and $B(S) = 10,000$ blocks. Assume all memory is used in the join process.

1. [6 points] Describe the steps of the algorithm (i.e., how many passes, what are the outputs generated at each pass, and the size of the outputs).

1st pass: hash R into 100 buckets, each bucket has a size of 50. Hash S using same hashing function into 100 buckets, each bucket has a size of 100.

2nd pass: read a entire bucket of S into memory one by one, and read a block of R 's data in the corresponding bucket one by one, join and output the join result.

2. [2 points] What is the total cost (i.e., the number of block I/O's) of the algorithm?

$$3B(R) + 3B(S) = 45000$$

3. [2 points] If the same join is performed using the **sort-merge** join algorithm, can it be done in **two** passes? Explain your answer.

No it can not.

Because $B(R) + B(S) > M^2$